

```
In [ ]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

## DATA COLLECTION AND PREPARATION :-

### IMPORTING THE REQUIRED LIBRARIES :-

```
In [ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import tensorflow as tf
from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import OrdinalEncoder
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler
from sklearn.compose import ColumnTransformer
from sklearn.model_selection import train_test_split
from sklearn.model_selection import RandomizedSearchCV, GridSearchCV
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
import tensorflow
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import pickle
import warnings
warnings.filterwarnings('ignore')
```

### COLLECT AND READ THE DATASET :-

```
In [ ]: df = pd.read_csv('/content/drive/MyDrive/Colab Notebooks/flightdata.csv')
```

```
In [ ]: pd.set_option('display.max_rows',100)
pd.set_option('display.max_columns',1000)
pd.set_option('display.width',1000)
```

```
In [ ]: df.head()
```

```
Out[ ]:
```

	YEAR	QUARTER	MONTH	DAY_OF_MONTH	DAY_OF_WEEK	UNIQUE_CARRIER	TAIL_NUM
0	2016	1	1	1	5	DL	N836I
1	2016	1	1	1	5	DL	N964I
2	2016	1	1	1	5	DL	N813I
3	2016	1	1	1	5	DL	N587N
4	2016	1	1	1	5	DL	N836I

## DESCRIPTIVE STATISTICAL :-

```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 11231 entries, 0 to 11230
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   YEAR                  11231 non-null  int64
1   QUARTER                11231 non-null  int64
2   MONTH                 11231 non-null  int64
3   DAY_OF_MONTH           11231 non-null  int64
4   DAY_OF_WEEK            11231 non-null  int64
5   UNIQUE_CARRIER        11231 non-null  object
6   TAIL_NUM               11231 non-null  object
7   FL_NUM                 11231 non-null  int64
8   ORIGIN_AIRPORT_ID      11231 non-null  int64
9   ORIGIN                 11231 non-null  object
10  DEST_AIRPORT_ID        11231 non-null  int64
11  DEST                   11231 non-null  object
12  CRS_DEP_TIME           11231 non-null  int64
13  DEP_TIME               11124 non-null  float64
14  DEP_DELAY              11124 non-null  float64
15  DEP_DEL15              11124 non-null  float64
16  CRS_ARR_TIME           11231 non-null  int64
17  ARR_TIME               11116 non-null  float64
18  ARR_DELAY              11043 non-null  float64
19  ARR_DEL15              11043 non-null  float64
20  CANCELLED              11231 non-null  float64
21  DIVERTED               11231 non-null  float64
22  CRS_ELAPSED_TIME       11231 non-null  float64
23  ACTUAL_ELAPSED_TIME    11043 non-null  float64
24  DISTANCE               11231 non-null  float64
25  Unnamed: 25            0 non-null      float64
dtypes: float64(12), int64(10), object(4)
memory usage: 2.2+ MB
```

```
In [ ]: df.isnull().sum()
```

```
Out[ ]: YEAR          0
        QUARTER       0
        MONTH         0
        DAY_OF_MONTH  0
        DAY_OF_WEEK   0
        UNIQUE_CARRIER 0
        TAIL_NUM       0
        FL_NUM         0
        ORIGIN_AIRPORT_ID 0
        ORIGIN         0
        DEST_AIRPORT_ID 0
        DEST           0
        CRS_DEP_TIME    0
        DEP_TIME       107
        DEP_DELAY       107
        DEP_DEL15       107
        CRS_ARR_TIME    0
        ARR_TIME       115
        ARR_DELAY       188
        ARR_DEL15       188
        CANCELLED       0
        DIVERTED        0
        CRS_ELAPSED_TIME 0
        ACTUAL_ELAPSED_TIME 188
        DISTANCE        0
        Unnamed: 25     11231
        dtype: int64
```

```
In [ ]: df.describe()
```

```
Out[ ]:
```

	YEAR	QUARTER	MONTH	DAY_OF_MONTH	DAY_OF_WEEK	FL_NUM
<b>count</b>	11231.0	11231.000000	11231.000000	11231.000000	11231.000000	11231.000000
<b>mean</b>	2016.0	2.544475	6.628973	15.790758	3.960199	1334.325617
<b>std</b>	0.0	1.090701	3.354678	8.782056	1.995257	811.875227
<b>min</b>	2016.0	1.000000	1.000000	1.000000	1.000000	7.000000
<b>25%</b>	2016.0	2.000000	4.000000	8.000000	2.000000	624.000000
<b>50%</b>	2016.0	3.000000	7.000000	16.000000	4.000000	1267.000000
<b>75%</b>	2016.0	3.000000	9.000000	23.000000	6.000000	2032.000000
<b>max</b>	2016.0	4.000000	12.000000	31.000000	7.000000	2853.000000

So , At last the columns that can be useful for prediction are...

```
In [ ]: df = df[['FL_NUM', 'MONTH', 'DAY_OF_MONTH', 'DAY_OF_WEEK', 'ORIGIN', 'DEST', 'C
```

```
In [ ]: df.isnull().sum()
```

```
Out[ ]: FL_NUM          0
        MONTH          0
        DAY_OF_MONTH   0
        DAY_OF_WEEK    0
        ORIGIN         0
        DEST           0
        CRS_ARR_TIME   0
        DEP_DEL15      107
        ARR_DEL15      188
        DEP_DELAY       107
        dtype: int64
```

HANDLING MISSING VALUES :-

```
In [ ]: df['DEP_DEL15'].mode()
```

```
Out[ ]: 0    0.0
        Name: DEP_DEL15, dtype: float64
```

```
In [ ]: df['ARR_DEL15'].mode()
```

```
Out[ ]: 0    0.0
        Name: ARR_DEL15, dtype: float64
```

```
In [ ]: df['DEP_DEL15'].fillna(0.0,inplace=True)
        df['ARR_DEL15'].fillna(0.0,inplace=True)
        df['DEP_DELAY'].fillna(df['DEP_DELAY'].median(),inplace=True)
```

<ipython-input-64-148c2a153b28>:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame  
  
See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
df['DEP\_DEL15'].fillna(0.0,inplace=True)

```
In [ ]: df.isnull().sum()
```

```
Out[ ]: FL_NUM          0
        MONTH          0
        DAY_OF_MONTH   0
        DAY_OF_WEEK    0
        ORIGIN         0
        DEST           0
        CRS_ARR_TIME   0
        DEP_DEL15       0
        ARR_DEL15       0
        DEP_DELAY       0
        dtype: int64
```

```
In [ ]: df.to_csv('df_reduced.csv')
        df.head()
```

```
Out[ ]:
```


	FL_NUM	MONTH	DAY_OF_MONTH	DAY_OF_WEEK	ORIGIN	DEST	CRS_ARR_TIME	DE
0	1399	1	1	5	ATL	SEA	2143	
1	1476	1	1	5	DTW	MSP	1435	
2	1597	1	1	5	ATL	SEA	1215	
3	1768	1	1	5	SEA	MSP	1335	
4	1823	1	1	5	SEA	DTW	607	

## EDA : EXPLORATORY DATA ANALYSIS :-

### UNIVARIATE ANALYSIS :-


```
In [ ]: sns.distplot(df['DEP_DELAY'],color='green',bins=20)
```

```
Out[ ]: <Axes: xlabel='DEP_DELAY', ylabel='Density'>
```




```
In [ ]: sns.histplot(df['CRS_ARR_TIME'],color='green',bins=20)
```

```
Out[ ]: <Axes: xlabel='CRS_ARR_TIME', ylabel='Count'>
```




```
In [ ]: sns.boxplot(df['CRS_ARR_TIME'])
```

```
Out[ ]: <Axes: >
```




```
In [ ]: sns.boxplot(df,y='DEP_DELAY')
```

```
Out[ ]: <Axes: ylabel='DEP_DELAY'>
```



```
In [ ]: plt.title('ARR_DEL15 : delayed at arrival more than 15 minutes')
plt.pie(df.ARR_DEL15.value_counts(), labels = ['delayed\nat arrival' if x
plt.show()
```


ARR\_DEL15 : delayed at arrival more than 15 minutes



BIVARIATE ANALYSIS :-

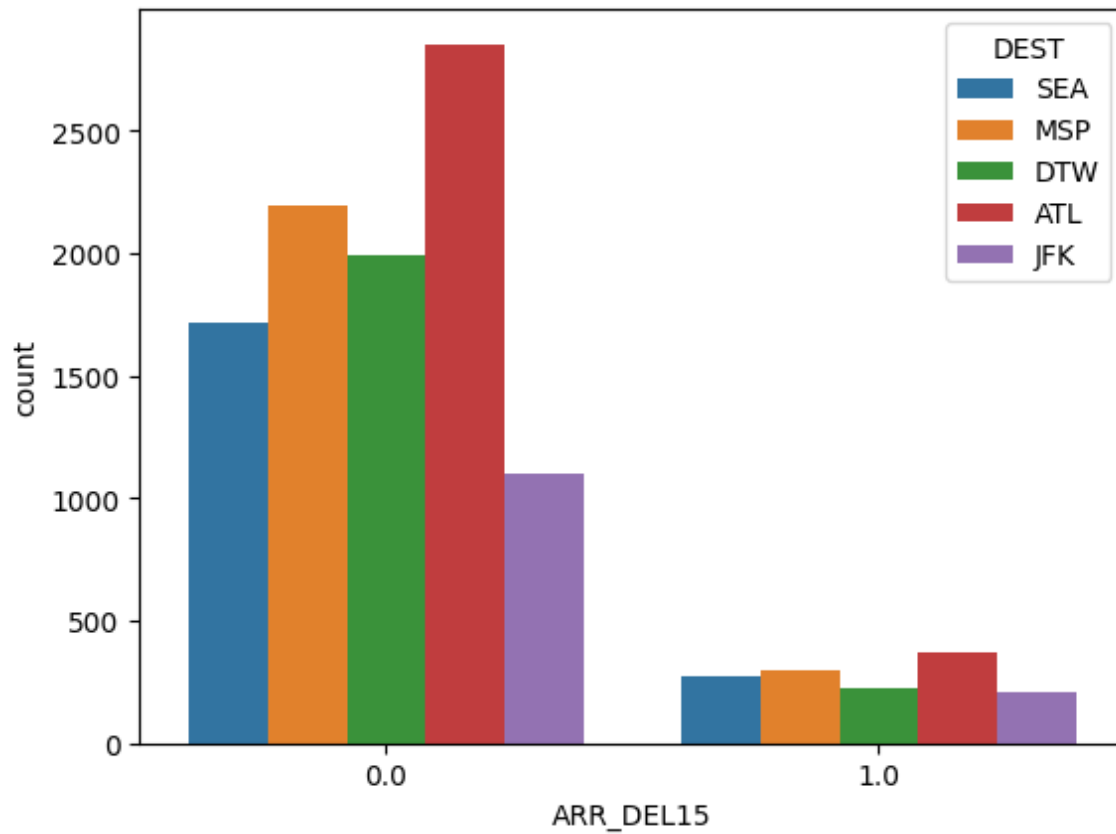
```
In [ ]: sns.countplot(data = df, x='ARR_DEL15', hue='MONTH')
```

```
Out[ ]: <Axes: xlabel='ARR_DEL15', ylabel='count'>
```



```
In [ ]: sns.countplot(data = df,x='ARR_DEL15',hue='DEST')
```

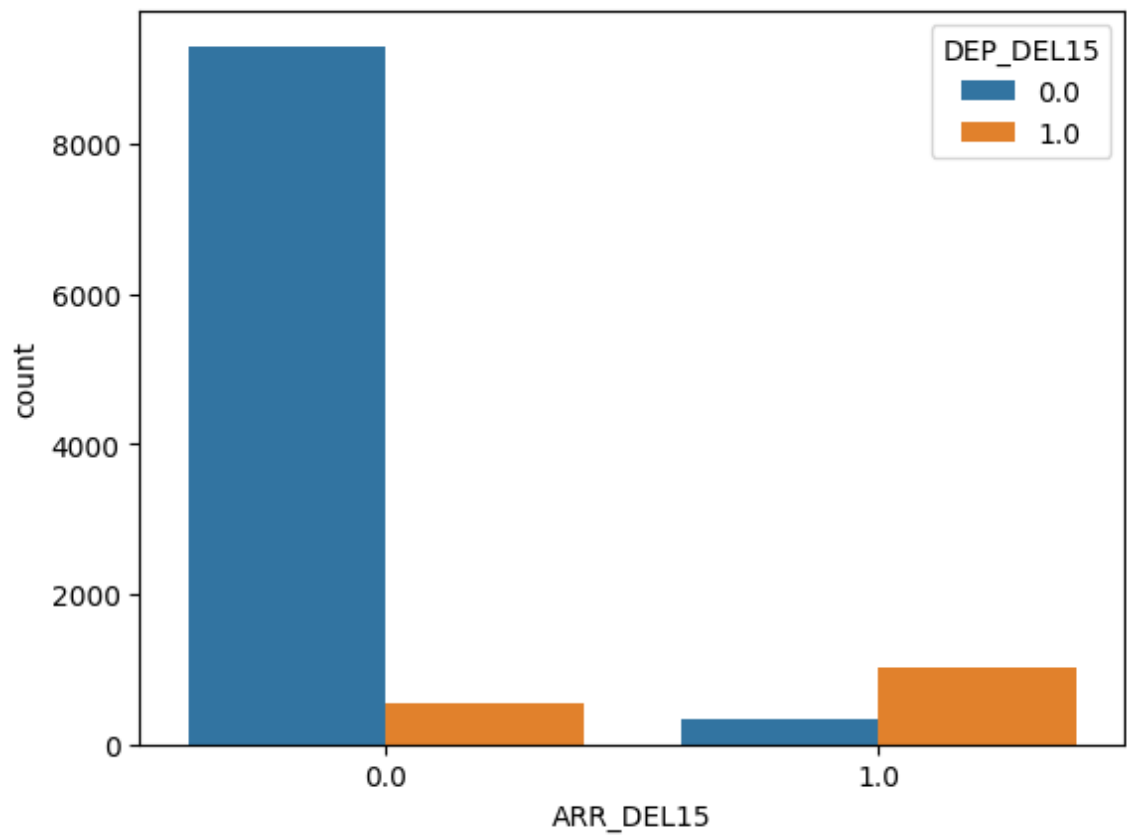
```
Out[ ]: <Axes: xlabel='ARR_DEL15', ylabel='count'>
```



```
In [ ]: sns.countplot(data = df,x='ARR_DEL15',hue='DEP_DEL15')
```


```
Out[ ]: <Axes: xlabel='ARR_DEL15', ylabel='count'>
```






```
In [ ]: sns.catplot(x='ARR_DEL15',y='DEP_DELAY',data=df)
```

```
Out[ ]: <seaborn.axisgrid.FacetGrid at 0x7f81cdc61a00>
```



```
In [ ]: sns.displot(df,x='ORIGIN',y='DEST')
```

```
Out[ ]: <seaborn.axisgrid.FacetGrid at 0x7f81c986bc40>
```




MULTIVARIATE ANALYSIS :-

```
In [ ]: sns.heatmap(df.corr(),annot=True)
```

<ipython-input-77-8df7bcac526d>:1: FutureWarning: The default value of numeric\_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric\_only to silence this warning.

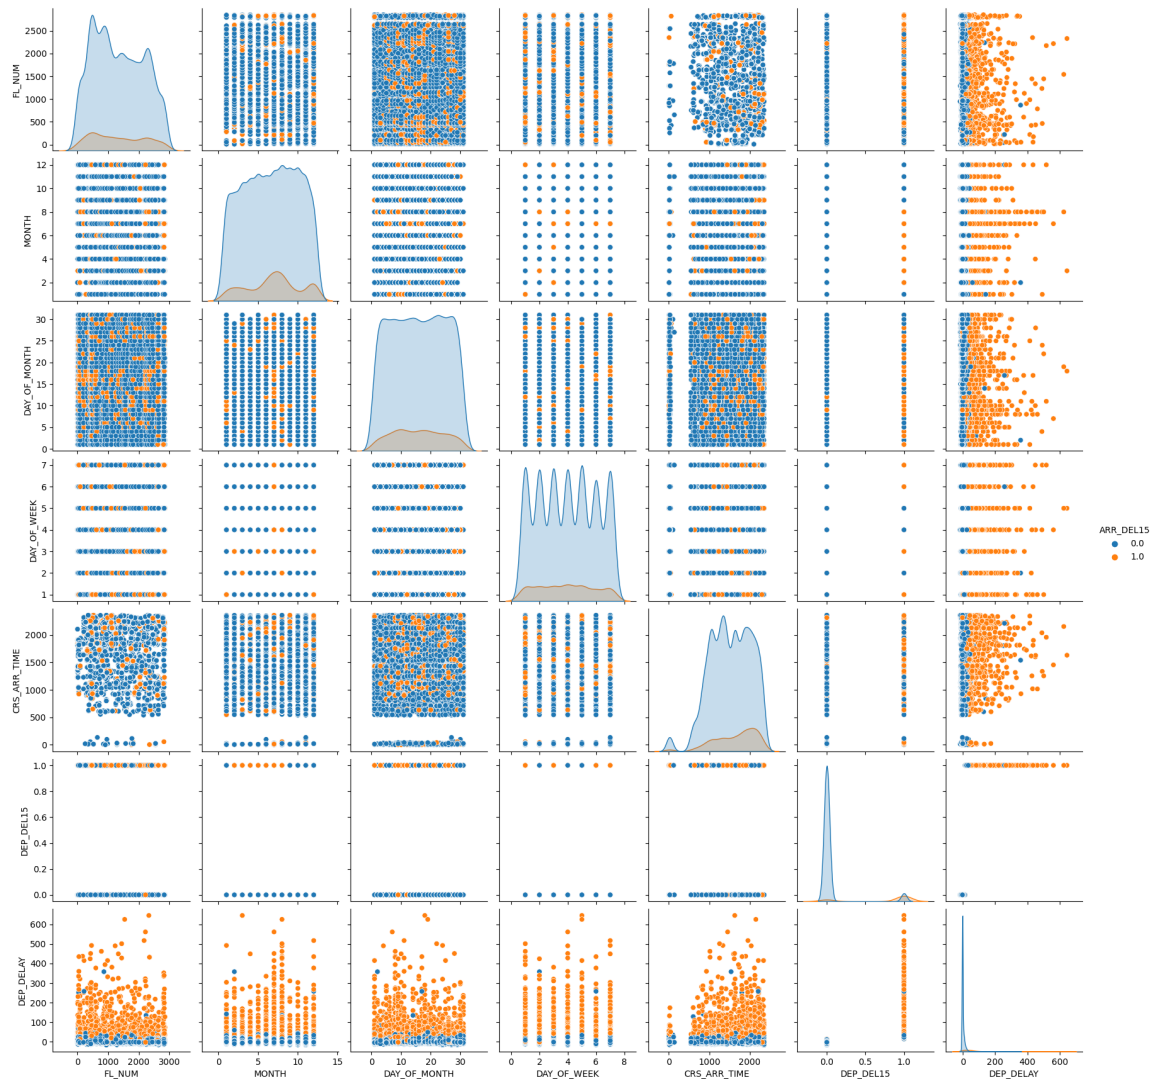
```
sns.heatmap(df.corr(),annot=True)
```

```
Out[ ]: <Axes: >
```



```
In [ ]: sns.pairplot(df, hue='ARR_DEL15')
```

```
Out[ ]: <seaborn.axisgrid.PairGrid at 0x7f81c8050fd0>
```



```
In [ ]: x = df[['FL_NUM', 'MONTH', 'DAY_OF_MONTH', 'DAY_OF_WEEK', 'ORIGIN', 'DEST', 'CRS_ARR_TIME', 'DEP_DELAY']]
y = df[['ARR_DEL15']]
```

HANDLING CATEGORICAL VALUES & SCALING THE DATA:-

```
In [ ]: ct1 = ColumnTransformer([('oe', OrdinalEncoder(), ['FL_NUM'])], ('ohe', OneHotEncoder(), ['ORIGIN', 'DEST']), remainder='passthrough')
ct2 = ColumnTransformer([('sc', StandardScaler(), ['CRS_ARR_TIME', 'DEP_DELAY'])], ('sc', StandardScaler(), ['CRS_ARR_TIME', 'DEP_DELAY']), remainder='passthrough')

x = pd.DataFrame(ct1.fit_transform(x), columns=ct1.get_feature_names_out())
x = pd.DataFrame(ct2.fit_transform(x), columns=ct2.get_feature_names_out())

x.head()
```

```
Out[ ]:   sc_oe_FL_NUM  sc_remainder_CRS_ARR_TIME  sc_remainder_DEP_DELAY  remainder_DEP_DELAY
0         0.109290                1.205371                -0.174060                0.0
1         0.239959                -0.203612                -0.256033                0.0
2         0.395756                -0.641431                -0.174060                0.0
3         0.581707                -0.402620                -0.201385                0.0
4         0.642016                -1.851405                -0.338007                0.0
```

```
In [ ]: pickle.dump(ct1,open('col_trans1.pkl','wb'))
        pickle.dump(ct2,open('col_trans2.pkl','wb'))
```

```
In [ ]: x.head()
```

```
Out[ ]:
```

	sc_oe_FL_NUM	sc_remainder_CRS_ARR_TIME	sc_remainder_DEP_DELAY	remainde
0	0.109290	1.205371	-0.174060	
1	0.239959	-0.203612	-0.256033	
2	0.395756	-0.641431	-0.174060	
3	0.581707	-0.402620	-0.201385	
4	0.642016	-1.851405	-0.338007	

```
In [ ]: y.head()
```

```
Out[ ]: 0    0.0
        1    0.0
        2    0.0
        3    0.0
        4    0.0
        Name: ARR_DEL15, dtype: float64
```

SPLITTING THE DATASET INTO TRAINING AND TESTING :-

```
In [ ]: x_train , x_test , y_train , y_test = train_test_split(x,y,test_size=0.2)
```

**MODEL BUILDING :-**

RANDOM FOREST MODEL :-

```
In [ ]: rfc = RandomForestClassifier()
        rfc.fit(x_train,y_train)
        y_pred = rfc.predict(x_test)

        acc = accuracy_score(y_test,y_pred)
        acc
```

```
Out[ ]: 0.9425901201602136
```

DECISION TREE MODEL :-

```
In [ ]: dtc = DecisionTreeClassifier()
        dtc.fit(x_train,y_train)
        y_pred = dtc.predict(x_test)

        acc = accuracy_score(y_test,y_pred)
        acc
```

```
Out[ ]: 0.9020916777926123
```

ANN MODEL :-

```
In [ ]: ann = Sequential()
ann.add(Dense(8,activation='relu'))
ann.add(Dense(32,activation='relu'))
ann.add(Dense(32,activation='relu'))
ann.add(Dense(1,activation='sigmoid'))

ann.compile(optimizer='adam',loss='binary_crossentropy',metrics=['accuracy'])

ann.fit(x_train,y_train,batch_size=4,validation_split=0.2,epochs=15)
```

```
Epoch 1/15
1797/1797 [=====] - 5s 2ms/step - loss: 0.2333
- accuracy: 0.9217 - val_loss: 0.2344 - val_accuracy: 0.9238
Epoch 2/15
1797/1797 [=====] - 4s 2ms/step - loss: 0.1771
- accuracy: 0.9421 - val_loss: 0.2213 - val_accuracy: 0.9238
Epoch 3/15
1797/1797 [=====] - 6s 3ms/step - loss: 0.1716
- accuracy: 0.9431 - val_loss: 0.2309 - val_accuracy: 0.9226
Epoch 4/15
1797/1797 [=====] - 4s 2ms/step - loss: 0.1679
- accuracy: 0.9438 - val_loss: 0.2247 - val_accuracy: 0.9277
Epoch 5/15
1797/1797 [=====] - 4s 2ms/step - loss: 0.1645
- accuracy: 0.9450 - val_loss: 0.2272 - val_accuracy: 0.9243
Epoch 6/15
1797/1797 [=====] - 6s 3ms/step - loss: 0.1625
- accuracy: 0.9467 - val_loss: 0.2304 - val_accuracy: 0.9260
Epoch 7/15
1797/1797 [=====] - 4s 2ms/step - loss: 0.1596
- accuracy: 0.9467 - val_loss: 0.2323 - val_accuracy: 0.9226
Epoch 8/15
1797/1797 [=====] - 4s 2ms/step - loss: 0.1557
- accuracy: 0.9489 - val_loss: 0.2344 - val_accuracy: 0.9221
Epoch 9/15
1797/1797 [=====] - 8s 4ms/step - loss: 0.1535
- accuracy: 0.9475 - val_loss: 0.2386 - val_accuracy: 0.9249
Epoch 10/15
1797/1797 [=====] - 4s 2ms/step - loss: 0.1498
- accuracy: 0.9521 - val_loss: 0.2380 - val_accuracy: 0.9243
Epoch 11/15
1797/1797 [=====] - 4s 2ms/step - loss: 0.1470
- accuracy: 0.9516 - val_loss: 0.2423 - val_accuracy: 0.9226
Epoch 12/15
1797/1797 [=====] - 5s 3ms/step - loss: 0.1451
- accuracy: 0.9524 - val_loss: 0.2501 - val_accuracy: 0.9226
Epoch 13/15
1797/1797 [=====] - 4s 2ms/step - loss: 0.1424
- accuracy: 0.9537 - val_loss: 0.2503 - val_accuracy: 0.9215
Epoch 14/15
1797/1797 [=====] - 4s 2ms/step - loss: 0.1399
- accuracy: 0.9527 - val_loss: 0.2648 - val_accuracy: 0.9182
Epoch 15/15
1797/1797 [=====] - 5s 3ms/step - loss: 0.1361
- accuracy: 0.9556 - val_loss: 0.2646 - val_accuracy: 0.9204
```

```
Out[ ]: <keras.callbacks.History at 0x7f81c02ac0a0>
```

```
In [ ]: y_pred = ann.predict(x_train)
```

```

y_pred = [0 if x<0.5 else 1 for x in y_pred]
acc = accuracy_score(y_train,y_pred)
print('train data prediction accuracy : ',acc)

y_pred = ann.predict(x_test)

y_pred = [0 if x<0.5 else 1 for x in y_pred]
acc = accuracy_score(y_test,y_pred)
print('test data prediction accuracy : ',acc)

```

```

281/281 [=====] - 1s 2ms/step
train data prediction accuracy : 0.9491317898486198
71/71 [=====] - 0s 1ms/step
test data prediction accuracy : 0.9345794392523364

```

HYPER PARAMETER TUNING :-

```

In [ ]: from scipy.stats import randint
params = {
    'n_estimators':[int(x) for x in np.linspace(50,500,50)],
    'criterion':['gini','entropy'],
    'max_features':['sqrt','log2'],
    'max_depth':[None,5,10,15,20,25,30],
    'min_samples_split':[int(x) for x in np.linspace(2,20)],
    'min_samples_leaf':[int(x) for x in np.linspace(1,20)],
}
rscv = RandomizedSearchCV(estimator=RandomForestClassifier(),param_distri
rscv.fit(x_train,y_train)

y_pred = rscv.predict(x_test)
acc = accuracy_score(y_pred,y_test)
print('accuracy score : ',acc)
print(rscv.best_params_)

```

```

accuracy score : 0.9457053849577214
{'n_estimators': 114, 'min_samples_split': 7, 'min_samples_leaf': 3, 'ma
x_features': 'sqrt', 'max_depth': 20, 'criterion': 'entropy'}

```

```

In [ ]: rfc2 = RandomForestClassifier(n_estimators= 114, min_samples_split= 7, mi
rfc2.fit(x_train,y_train)
y_pred = rfc2.predict(x_test)

acc = accuracy_score(y_test,y_pred)
print('accuracy score : ',acc)

```

```

accuracy score : 0.945260347129506

```

```

In [ ]: params = {
    'max_depth':list(range(3,14,2)),
    'criterion':['gini','entropy'],
    'min_samples_split':list(range(2,11,2)),
    'min_samples_leaf':list(range(1,6))
}
gscv = GridSearchCV(estimator=DecisionTreeClassifier(),param_grid=params,
gscv.fit(x_train,y_train)

y_pred = gscv.predict(x_test)
acc = accuracy_score(y_pred,y_test)
print('accuracy score : ',acc)
print(gscv.best_params_)

```

```
accuracy score : 0.9434801958166444
{'criterion': 'entropy', 'max_depth': 5, 'min_samples_leaf': 1, 'min_samples_split': 2}
```

```
In [ ]: dtc2 = DecisionTreeClassifier(criterion= 'entropy', max_depth= 5, min_samples_split= 2)
dtc2.fit(x_train,y_train)
y_pred = dtc2.predict(x_test)

acc = accuracy_score(y_test,y_pred)
print('accuracy score : ',acc)
```

```
accuracy score : 0.943035157988429
```

TESTING THE MODEL WITH MULTIPLE EVALUATION METRICS (AFTER HYPER  
PARAMETER TUNING ):-

```
In [ ]: def cl_res(name,model):
    y_pred = model.predict(x_test)
    if(name=='artificial_neural_network'):
        y_pred = [0 if x<0.5 else 1 for x in y_pred]
    print(name,' :-\n-----')
    print('accuracy score of ',name,' : ',accuracy_score(y_test,y_pred))
    print(classification_report(y_test,y_pred,target_names=['no delay','delay']))
    print('confusion matrix : \n',confusion_matrix(y_test,y_pred))
    print('\n')
    # plt.subplot(121)
    plt.figure(figsize=(3,2))
    sns.heatmap(confusion_matrix(y_test,y_pred),annot=True)
    # plt.subplot(122)
    plt.figure(figsize=(1,1))
    RocCurveDisplay.from_predictions(y_test,y_pred)
    plt.show()
    print('\n\n')
```

```
In [ ]: cl_res('random_forest_classifier(before tuning)',rfc)
```

```
random_forest_classifier(before tuning) :-
```

```
-----
```


```
accuracy score of random_forest_classifier(before tuning) : 0.9425901201602136
```

	precision	recall	f1-score	support
no delay	0.95	0.98	0.97	1962
delay	0.84	0.68	0.75	285
accuracy			0.94	2247
macro avg	0.90	0.83	0.86	2247
weighted avg	0.94	0.94	0.94	2247


```
confusion matrix :
```

```
[[1924  38]
 [ 91 194]]
```





<Figure size 100x100 with 0 Axes>



```
In [ ]: cl_res('random_forest_classifier(after tuning)', rfc2)
```

```

random_forest_classifier(after tuning) :-
-----
accuracy score of random_forest_classifier(after tuning) : 0.94526034
7129506


```

	precision	recall	f1-score	support
no delay	0.96	0.98	0.97	1962
delay	0.85	0.69	0.76	285
accuracy			0.95	2247
macro avg	0.90	0.84	0.87	2247
weighted avg	0.94	0.95	0.94	2247


```

confusion matrix :
[[1928  34]
 [ 89 196]]

```



<Figure size 100x100 with 0 Axes>




```
In [ ]: cl_res('decision_tree_classifier(before tuning)',dtc)
```


```
decision_tree_classifier(before tuning) :-
-----
accuracy score of decision_tree_classifier(before tuning) : 0.9020916
777926123
```

	precision	recall	f1-score	support
no delay	0.95	0.93	0.94	1962
delay	0.60	0.69	0.64	285
accuracy			0.90	2247
macro avg	0.78	0.81	0.79	2247
weighted avg	0.91	0.90	0.91	2247

```
confusion matrix :
[[1829 133]
 [ 87 198]]
```



<Figure size 100x100 with 0 Axes>



```
In [ ]: cl_res('decision_tree_classifier(after tuning)',dtt2)
```

```
decision_tree_classifier(after tuning) :-
```


```
-----
```

```
accuracy score of decision_tree_classifier(after tuning) : 0.943035157988429
```


	precision	recall	f1-score	support
no delay	0.95	0.99	0.97	1962
delay	0.87	0.65	0.74	285
accuracy			0.94	2247
macro avg	0.91	0.82	0.85	2247
weighted avg	0.94	0.94	0.94	2247

```
confusion matrix :
```

```
[[1935  27]
 [ 101 184]]
```



<Figure size 100x100 with 0 Axes>



```
In [ ]: cl_res('artificial_neural_network',ann)
```

71/71 [=====] - 0s 1ms/step

artificial\_neural\_network :-

-----

accuracy score of artificial\_neural\_network : 0.9345794392523364

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

no delay	0.95	0.97	0.96	1962
----------	------	------	------	------

delay	0.78	0.68	0.73	285
-------	------	------	------	-----

accuracy			0.93	2247
----------	--	--	------	------


macro avg	0.87	0.83	0.84	2247
-----------	------	------	------	------

weighted avg	0.93	0.93	0.93	2247
--------------	------	------	------	------


confusion matrix :

[[1906 56]

[ 91 194]]



<Figure size 100x100 with 0 Axes>



- Here Random Forest Classifier (after tuning) has the highest accuracy score and good at other evaluation metrics, so we are going to save that model.

#### SAVING THE MODEL :-

```
In [ ]: pickle.dump(rfc2,open('random_forest_classifier.pkl','wb'))
```